

AMQ Streams command cheat sheet

[Introduction](#)

[Adjust environment variables](#)

[Kafka Cluster and Kafka Broker](#)

[Kafka cluster definition](#)

[Broker configurations](#)

[Show cluster-wide finalized features and features supported by a specific broker](#)

[Show the controller Broker ID](#)

[Change the controller broker to others](#)

[Living Kafka broker list using Zookeeper](#)

[Kafka broker information on Zookeeper using Zookeeper](#)

[Delete and recreate a Kafka pod and its PVC](#)

[Manual rolling update: Kafka brokers](#)

[Zookeeper](#)

[Zookeeper definition](#)

[Zookeeper interactive command line](#)

[Status and environment of Zookeeper](#)

[Dump sessions of zookeeper](#)

[Manual rolling update: Zookeeper](#)

[Access Zookeepers directly for debugging : since AMQ Streams 1.6](#)

[Create a zoo-entrance to access Zookeepers directly for debugging : until AMQ Streams 1.5](#)

[Show internal zookeeper snapshot and zookeeper log](#)

[Kafka Topic](#)

[Kafka topic definitions](#)

[Actual Kafka topic list](#)

[Actual Kafka topic description](#)

[Kafka topic configurations](#)

[Create a Kafka topic \(replication factor: 3, partitions: 5\)](#)

[Delete a Kafka topic](#)

[Topic latest offset \(for each partition\)](#)

[Topic earliest offset \(for each partition\)](#)

[Topic size \(based on offsets of total partitions\)](#)

[Creation timestamp of a Kafka topic](#)

[Disable topic deletion](#)

[Kafka Internal Topic](#)

[Show consumer offsets contents](#)

[Show transaction state contents](#)

[Strimzi Internal Topic \(since AMQ Streams 1.7\)](#)

[Show strimzi_store_topic](#)

[Show strimzi-topic-operator-kstreams-topic-store-changelog](#)

[Kafka Partition and Replica](#)

[Log directories and partitions](#)

[Reassign partitions : since AMQ Streams 1.6](#)

[Reassign partitions : until AMQ Streams 1.5](#)

[Stop Reassigning partitions](#)

[Topics that have under-replicated-partitions](#)

[Topics that have unavailable-partitions](#)

[Preferred replica election](#)

[Kafka Producer](#)

[Console producer](#)

[Producer performance test tool \(size:10, throughput:1, records:1000\)](#)

[Kafka Consumer and Consumer Group](#)

- [Console consumer \(without a consumer group\)](#)
- [Console consumer \(with a consumer group\)](#)
- [Consumer performance test tool](#)
- [Consumer group list](#)
- [Consumer group description \(--offset\)](#)
- [Consumer group description \(--state\)](#)
- [Delete a consumer group](#)
- [Reset the offsets for a specific consumer group](#)
- [Delete offset with specific topic for a specific consumer group](#)

[About Kafka CLI execution environment](#)

- [Executing Kafka CLI using an existing broker](#)
- [Executing Kafka CLI using a new pod \(do not use a pod of an existing broker\)](#)
- [Executing Kafka CLI using a customized new pod as root](#)

[Strimzi Cluster Operator](#)

- [Stop cluster operator](#)
- [Pause reconciliation of custom resources](#)
- [Check for duplicate cluster operators installed](#)

[Strimzi Entity Operator \(comprises Topic Operator and User Operator\)](#)

- [Stop entity operator](#)

[Kafka Connect and Debezium](#)

- [Create a MySQL DB POD for test](#)
- [Create Kafka Connect image with Debezium: build a image on the OpenShift](#)
- [Create a MySQL Debezium connector accessing to MySQL DB on the Kafka Connector POD](#)

[Cruise Control](#)

- [Cruise Control definition](#)

[TLS](#)

- [Validity days of Cluster CA and Clients CA certifications in secrets](#)
- [Check actual TLS info using openssl s_client command](#)
- [Check TLS port on a Kafka broker using openssl s_client command](#)
- [Enable ssl debug log on Kafka brokers](#)
- [Renewing CA certificates manually](#)
- [Replace to my own cluster ca certification and its key](#)
- [Enable SSL debug log of Kafka CLI for debugging](#)
- [Cluster CA and truststore.p12](#)
- [KafkaUser secret and keystore.p12](#)

[tls listener \(port:9093\) with one way TLS](#)

- [Prerequisites](#)
- [Console producer](#)
- [Producer performance test tool \(size:10, throughput:1, records:1000\)](#)
- [Console consumer](#)

[tls listener \(port:9093\) with two-way\(mutual\) TLS client authentication and simple authorization](#)

- [Prerequisites](#)
- [Console producer](#)
- [Producer performance test tool \(size:10, throughput:1, records:1000\)](#)
- [Console consumer \(with a consumer group\)](#)

[JMX](#)

- [Enable JMX port\(9999\) on the Kafka brokers](#)
- [Print Kafka metrics via JMX using kafka.tools.JmxTool](#)
- [jmxterm](#)
- [jconsole \(or jmxterm\) from localhost machine to a Kafka broker](#)
- [jmxterm from a new pod to a Kafka broker](#)

[JMX Exporter and Kafka Exporter](#)

- [Show JMX exporter metrics using curl](#)
- [Show Kafka exporter metrics using curl](#)

Log

[Log of a pod](#)

[Multiple pod logs using stern](#)

[Enable debug log : Kafka, Zookeeper, Entity Operator, Cruise Control and its tls-sidecar](#)

[Enable debug log : Cluster Operator](#)

[Enable debug log : Kafka Bridge](#)

[Enable debug log : Kafka Connect](#)

[Enable debug log : Kafka Mirror Maker v1](#)

[Enable debug log : Kafka Mirror Maker v2](#)

Kafka Client Example

[Run Kafka producer client example on the OpenShift](#)

[Run Kafka streams client example on the OpenShift](#)

Debug

[JavaVM Info](#)

[Thread Dump](#)

[Packet capture using tcpdump](#)

[Heap Histogram & Stastics](#)

[Heap Dump](#)

[RHEL8 support tool image](#)

[Change AMQ Streams Kafka image: build a image on the local machine](#)

[Change AMQ Streams Kafka image: build a image on the OpenShift](#)

[Remote Debugging: Kafka Broker](#)

[Remote Debugging: Cluster Operator](#)

[Remote Debugging: Topic Operator in Entity Operator](#)

[Remote Debugging: User Operator in Entity Operator](#)

[Remote Debugging: Kafka Connect](#)

[Remote Debugging: Cruise Control](#)

Tips

[Dump OpenShift project](#)

[OpenShift events \(sort by last timestamp\)](#)

[Force a pod termination once](#)

[Continues to force a pod termination](#)

[Login to Code Ready Container host node using ssh](#)

[no matches for kind "Deployment" in version "extensions/v1beta1" error occurs](#)

[Long and short names for each Strimzi resource](#)

Dump AMQ Streams project

[Dump resources, status, and logs of AMQ Streams on OpenShift](#)

[Customizable report script](#)

References

Unorganized

[Login to the pod's node using debug container](#)

[Execute GDB from a node](#)

[Kill a connection using GDB](#)

[Show actual network policy](#)

Introduction

This is a command cheat sheet for AMQ Streams to help support engineers with their work.

Basically, most of the commands were tested with AMQ Streams 1.5, but they may not work with different versions of AMQ Streams. If you find a command that does not work, please leave comments with your environment info. it may be fixed later.

You are free to add any useful commands to this cheat sheet. Or, if you leave comments about a command you want, it may be added later.

Adjust environment variables

```
# Kafka cluster name
KAFKA_CLUSTER=my-cluster
# Kafka topic name
TOPIC_NAME=my-topic
# consumer group name
CONSUMER_GROUP_NAME=my-group
# Kafka user name
KAFKAUSER_NAME=my-user
# Kafka connect cluster
KAFKACONNECT_CLUSTER=my-connect-cluster
# Kafka broker ID
BROKER_ID=0
# Zookeeper ID
ZOOKEEPER_ID=0
# namespace name (i.e. project name)
NAMESPACE=my-project
```

Adding these adjusted environment variables makes it easier to run the following commands.

Kafka Cluster and Kafka Broker

Kafka cluster definition

```
oc get Kafka -o yaml
oc get statefulsets ${KAFKA_CLUSTER}-kafka -o yaml

# internal Kafka configuration
oc exec ${KAFKA_CLUSTER}-kafka-0 -- cat /tmp/strimzi.properties
oc exec ${KAFKA_CLUSTER}-kafka-0 -- cat /opt/kafka/custom-config/log4j.properties
oc exec ${KAFKA_CLUSTER}-kafka-0 -- ps aux
```

Broker configurations

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-configs.sh --bootstrap-server
localhost:9092 --entity-type brokers --entity-name ${BROKER_ID} --describe --all
```

Show cluster-wide finalized features and features supported by a specific broker

```
# cluster-wide finalized features
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-features.sh --bootstrap-server
localhost:9092 --describe

# features supported by a specific broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-features.sh --bootstrap-server
localhost:9092 --describe --
```

For details, refer to [KIP-584](#).

Show the controller Broker ID

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:12181 get
/controller
```

```
# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 get /controller
```

Change the controller broker to others

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 delete /controller

# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 delete /controller
```

✕cannot know which broker will be the next controller broker

Living Kafka broker list using Zookeeper

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 ls /brokers/ids

# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 ls /brokers/ids
```

Kafka broker information on Zookeeper using Zookeeper

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 get /brokers/ids/${BROKER_ID}

# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 get /brokers/ids/${BROKER_ID}
```

Delete and recreate a Kafka pod and its PVC

```
oc annotate pod ${KAFKA_CLUSTER}-kafka-${BROKER_ID} strimzi.io/delete-pod-and-pvc=true
```

Manual rolling update: Kafka brokers

```
oc annotate statefulset ${KAFKA_CLUSTER}-kafka strimzi.io/manual-rolling-update=true

# FYI as another way to restart, using patch adding date label to change CR
oc patch statefulset ${KAFKA_CLUSTER}-kafka -p
'{"spec":{"template":{"metadata":{"labels":{"date":"'`date +%s'`"}}}}}'
```

Wait for the next reconciliation to occur (every two minutes by default).

Zookeeper

Zookeeper definition

```
oc get Kafka -o yaml
oc get statefulsets ${KAFKA_CLUSTER}-zookeeper -o yaml

# internal Zookeeper configuration
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- cat /tmp/zookeeper.properties
```

```
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- cat /opt/kafka/custom-config/log4j.properties
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- ps aux
```

Zookeeper interactive command line

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:12181

# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181
```

Status and environment of Zookeeper

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bash -c "echo stat | nc localhost 12181"
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bash -c "echo envi | nc localhost 12181"

# until AMQ Stream 1.5 : cannot know which Zookeeper [0 - X] to connect to
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- bash -c "echo stat | nc localhost 2181"
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- bash -c "echo envi | nc localhost 2181"
```

Dump sessions of zookeeper

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bash -c "echo dump | nc localhost 12181"

# until AMQ Stream 1.5 : cannot know which Zookeeper [0 - X] to connect to
oc exec ${KAFKA_CLUSTER}-kafka-0 -- bash -c "echo dump | nc localhost 2181"
```

You can confirm the controller broker and living Kafka brokers

Manual rolling update: Zookeeper

```
oc annotate statefulset ${KAFKA_CLUSTER}-zookeeper strimzi.io/manual-rolling-update=true

# FYI as another way to restart Zookeepers, using patch adding date label to change CR
oc patch statefulset ${KAFKA_CLUSTER}-zookeeper -p
"{\"spec\":{\"template\":{\"metadata\":{\"labels\":{\"date\":\"`date +%s`\"}}}}}"
```

Wait for the next reconciliation to occur (every two minutes by default).

Access Zookeepers directly for debugging : since AMQ Streams 1.6

Since AMQ Streams 1.6, you don't need to create zoo-entrance to access Zookeepers directly. Zookeepers expose 12181 port as plain in the pods. So you can use port-forward.

```
# Step 1: create forward port to 12181 on the zookeeper
oc port-forward ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} 12181:12181

# Step 2: you can execute zookeeper-shell.sh from local machine via forward port
./bin/zookeeper-shell.sh 127.0.0.1:12181 get /controller
```

Create a zoo-entrance to access Zookeepers directly for debugging : until AMQ Streams 1.5

When Strimzi deploys the Kafka cluster, it keeps the ZooKeeper secure and inaccessible from any other application. This deploys a Stunnel proxy which exposes the Zookeeper without any authentication and encryption. For details, refer to <https://github.com/scholzj/zoo-entrance>.

```
# get zoo-entrance definition from the zoo-entrance project
wget -O ./zoo-entrance.yaml https://raw.githubusercontent.com/scholzj/zoo-entrance/master/deploy.yaml
# change Kafka cluster name
sed -i "s|my-cluster|${KAFKA_CLUSTER}|g" zoo-entrance.yaml
```

```
# create the deployment with the proxy named zoo-entrance, a service named zoo-entrance and a
Network Policy to allow the proxy to connect to the broker.
oc apply -f ./zoo-entrance.yaml

# test zoo-entrance
#   access from a broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- bin/zookeeper-shell.sh zoo-entrance:2181 ls
/brokers/ids
#   access from a new pod
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- bin/zookeeper-shell.sh zoo-entrance:2181 ls /brokers/ids

# [if you need] define NodePort to access Zookeeper from external
oc expose deployment/zoo-entrance --type=NodePort --name=zoo-entrance-external --port 2181
# access from external
./bin/zookeeper-shell.sh $(oc get ip):$(oc get service zoo-entrance-external
-o=jsonpath='{.spec.ports[0].nodePort}')
```

Show internal zookeeper snapshot and zookeeper log

```
# show zookeeper snapshot
oc exec -it ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -- env - java -cp './libs/*'
org.apache.zookeeper.server.SnapshotFormatter /var/lib/zookeeper/data/version-2/snapshot.<snapshot
number>

# show zookeeper log
oc exec -it ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -- env - java -cp './libs/*'
org.apache.zookeeper.server.LogFormatter /var/lib/zookeeper/data/version-2/log.<log number>
```

Kafka Topic

Kafka topic definitions

```
oc get KafkaTopic -o yaml
```

If you use a name that cannot be used for the resource name like underbar of “test_topic02” Kafka topic, the hash value will be added as a postfix as below:

```
oc get KafkaTopic
~~~
NAME                                PARTITIONS  REPLICATION FACTOR
testtopic01                          5            3
test-topic02---d04f2b34bd744b57f59e223c84cfc6aedef8b5391  5            3
~~~
```

Actual Kafka topic list

```
# using a broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --list
~~~
my-topic-1
my-topic-2
~~~

# using a new pod
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
```

```
--restart=Never -- bin/kafka-topics.sh --bootstrap-server ${KAFKA_CLUSTER}-kafka-bootstrap:9092
--list
```

Actual Kafka topic description

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --describe ${TOPIC_NAME}

~~~
Topic: my-topic-1      PartitionCount: 3      ReplicationFactor: 3      Configs:
segment.bytes=1073741824,retention.ms=7200000,message.format.version=2.5-IV0
  Topic: my-topic-1      Partition: 0      Leader: 2      Replicas: 2,1,0 Isr: 0,1,2
  Topic: my-topic-1      Partition: 1      Leader: 1      Replicas: 1,0,2 Isr: 0,1,2
  Topic: my-topic-1      Partition: 2      Leader: 0      Replicas: 0,2,1 Isr: 0,1,2
~~~
```

The ID of the first broker in “Replicas” is the preferred replica. If this is different from the “Leader”, then the cluster is out of balance, [kafka-preferred-replica-election.sh](#) may need to be used.

“Isr” means in-sync-replicas. If “Isr” is different from “Replicas”, it indicates that replication is delayed for some reason.

Kafka topic configurations

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-configs.sh --bootstrap-server
localhost:9092 --entity-type topics --entity-name ${TOPIC_NAME} --describe --all
```

Create a Kafka topic (replication factor: 3, partitions: 5)

```
# using kafka-topics.sh command
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --create --replication-factor 3 --partitions 5 --topic ${TOPIC_NAME}

# using CR
oc apply -f examples/topic/kafka-topic.yaml
```

The topic creation is done asynchronously.

Delete a Kafka topic

```
# using Kafka command
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --delete --topic ${TOPIC_NAME}

# using CR
oc delete KafkaTopic ${TOPIC_NAME}
```

The [topic deletion](#) must be enabled.

The topic deletion is done asynchronously. To verify that the topic has been deleted, please refer to [How to confirm that the topic has been deleted](#).

Topic latest offset (for each partition)

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-run-class.sh
kafka.tools.GetOffsetShell --broker-list localhost:9092 --topic ${TOPIC_NAME}
```

Topic earliest offset (for each partition)

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-run-class.sh
kafka.tools.GetOffsetShell --broker-list localhost:9092 --topic ${TOPIC_NAME} --time -2
```


Topic size (based on offsets of total partitions)

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-run-class.sh
kafka.tools.GetOffsetShell --broker-list localhost:9092 --topic ${TOPIC_NAME} --time -1 --offsets 1
| awk -F ":" '{sum += $3} END {print sum}'
```

Creation timestamp of a Kafka topic

```
# using zookeeper's ctime
# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 stat
/brokers/topics/${TOPIC_NAME}

# using creationTimestamp of Kafka topic CR
oc get KafkaTopic/${TOPIC_NAME} -o yaml -o=jsonpath='{.metadata.creationTimestamp}'
```

Disable topic deletion

```
kind: Kafka
...
spec:
  kafka:
    config:
      ...
      delete.topic.enable: false
```

With this setting, Kafka will no longer accept topic deletion instructions. Sometimes, we guide this procedure for additional safety. For example, if you suspect that the Topic Operator deletes topics during a reconciliation or upgrade.

Kafka Internal Topic

Show `__consumer_offsets` contents

```
oc run kafka-consumer-offsets -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
--rm=true --restart=Never -- bin/kafka-console-consumer.sh --formatter
"kafka.coordinator.group.GroupMetadataManager$OffsetsMessageFormatter" --bootstrap-server
${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --topic __consumer_offsets
```

Show `__transaction_state` contents

```
oc run kafka-transaction-state -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
--rm=true --restart=Never -- bin/kafka-console-consumer.sh --formatter
"kafka.coordinator.transaction.TransactionLog$TransactionLogMessageFormatter" --bootstrap-server
${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --topic __transaction_state
```

Producer's transaction id(transactional.id) is used to determine `__transaction_state` partition.

Strimzi Internal Topic (since AMQ Streams 1.7)

Show `__strimzi_store_topic`

```
oc run strimzi-store-topic -ti --image=registry.redhat.io/amq7/amq-streams-kafka-27-rhel7:1.7.0
--rm=true --restart=Never -- bin/kafka-console-consumer.sh --bootstrap-server
${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --topic __strimzi_store_topic
```

Show __strimzi-topic-operator-kstreams-topic-store-changelog

```
oc run strimzi-topic-operator-kstreams-topic-store-changelog -ti
--image=registry.redhat.io/amq7/amq-streams-kafka-27-rhel7:1.7.0 --rm=true --restart=Never --
bin/kafka-console-consumer.sh --bootstrap-server ${KAFKA_CLUSTER}-kafka-bootstrap:9092
--from-beginning --topic __strimzi-topic-operator-kstreams-topic-store-changelog
```

“__strimzi-topic-operator-kstreams-topic-store-changelog” is used as key value store in the internal Kafka Streams in the Topic Operator. For details, refer to [TopicStoreTopologyProvider](#) class.

Kafka Partition and Replica

Log directories and partitions

```
# using kafka-log-dirs.sh
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-log-dirs.sh --bootstrap-server
localhost:9092 --describe

~~~
Received log directory information from brokers 0,1,2
{"version":1,"brokers":[{"broker":0,"logDirs":[{"logDir":"/var/lib/kafka/data-0/kafka-log0","error":null,"partitions":[{"partition":"my-topic-1-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-0","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-0","size":0,"offsetLag":0,"isFuture":false}]}]},{"broker":1,"logDirs":[{"logDir":"/var/lib/kafka/data-0/kafka-log1","error":null,"partitions":[{"partition":"my-topic-1-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-0","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-0","size":0,"offsetLag":0,"isFuture":false}]}]},{"broker":2,"logDirs":[{"logDir":"/var/lib/kafka/data-0/kafka-log2","error":null,"partitions":[{"partition":"my-topic-1-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-0","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-1","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-1-2","size":0,"offsetLag":0,"isFuture":false}, {"partition":"my-topic-2-0","size":0,"offsetLag":0,"isFuture":false}]}]}]}

# FYI, actual log directories in a broker
# JBOD : "log.dirs" option in the kafka configuration
oc exec -it ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- find /var/lib/kafka/data-${JBOD_ID}/ #
not JBOD: "log.dir" option in the kafka configuration
oc exec -it ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- find /var/lib/kafka/data/
```

Reassign partitions : since AMQ Streams 1.6

```
# step 0:create forward port
oc port-forward ${KAFKA_CLUSTER}-zookeeper-0 12181:12181

# step 1: choose the partitions that should be reassigned
cat << "EOF" > topicPartitionList.json
{"topics": [{"topic": "my-topic"}],
"version":1
}
EOF

# step 2: export proposed partitions and current partitions
./bin/kafka-reassign-partitions.sh --zookeeper=localhost:12181 --generate --topics-to-move-json-file
topicPartitionList.json --broker-list 0,1,2,3 | tee >(tail -1 > proposed_partition.json) | tee
>(sed -n 2p > current_partition_`date +%Y%m%d_%H%M%S`.json)

# step 3: modify proposed_partition.json if you need

# step 4: reassign the partitions
```

```
./bin/kafka-reassign-partitions.sh --zookeeper=localhost:12181 --execute --reassignment-json-file
proposed_partition.json --throttle 5000000

# step 5: verify the partitions
# using --verify command
./bin/kafka-reassign-partitions.sh --zookeeper=localhost:12181 --verify --reassignment-json-file
proposed_partition.json
# using /admin/reassign_partitions znode
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- bin/zookeeper-shell.sh 127.0.0.1:2181 get
/admin/reassign_partitions
```

Reassign partitions : until AMQ Streams 1.5

```
# step 1: choose the partitions that should be reassigned
cat << "EOF" > topicPartitionList.json
{"topics": [{"topic": "my-topic"}],
  "version":1
}
EOF
# step 2: export proposed partitions and current partitions
./bin/kafka-reassign-partitions.sh --zookeeper $(crc ip):$(oc get service zoo-entrance-external
-o=jsonpath='{.spec.ports[0].nodePort}') --generate --topics-to-move-json-file
topicPartitionList.json --broker-list 0,1,2,3 | tee >(tail -1 > proposed_partition.json) | tee
>(sed -n 2p > current_partition_date "+%Y%m%d_%H%M%S".json)

# step 3: modify proposed_partition.json if you need

# step 4: reassign the partitions
./bin/kafka-reassign-partitions.sh --zookeeper $(crc ip):$(oc get service zoo-entrance-external
-o=jsonpath='{.spec.ports[0].nodePort}') --execute --reassignment-json-file proposed_partition.json
--throttle 5000000

# step 5: verify the partitions
# using --verify command
./bin/kafka-reassign-partitions.sh --zookeeper $(crc ip):$(oc get service zoo-entrance-external
-o=jsonpath='{.spec.ports[0].nodePort}') --verify --reassignment-json-file proposed_partition.json
# using /admin/reassign_partitions znode
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- bin/zookeeper-shell.sh 127.0.0.1:2181 get
/admin/reassign_partitions
```

For details, refer to the [document](#), and this example is using [zoo-entrance](#) to access Zookeeper.

Stop Reassigning partitions

```
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 delete
/admin/reassign_partitions

# until AMQ Stream 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 delete
/admin/reassign_partitions
```

In some circumstances, reassignment may be halted and re-execution may not be possible. In such cases, you can stop reassignment with this command. However, general use is not recommended.

Topics that have under-replicated-partitions

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --describe --under-replicated-partitions
```

Topics that have unavailable-partitions

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --unavailable-partitions
```

Preferred replica election

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-preferred-replica-election.sh --bootstrap-server localhost:9092
```

No major side effects, so if the cluster performance is not stable, run this command anyway.

Kafka Producer

Console producer

```
# using a new pod
oc run kafka-producer -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true --restart=Never -- bin/kafka-console-producer.sh --broker-list ${KAFKA_CLUSTER}-kafka-bootstrap:9092 --topic ${TOPIC_NAME}

# using kafka-0 broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -- env - bin/kafka-console-producer.sh --broker-list localhost:9092 --topic ${TOPIC_NAME}
```

Producer performance test tool (size:10, throughput:1, records:1000)

```
oc run kafka-producer-perf -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true --restart=Never -- bin/kafka-producer-perf-test.sh --producer-props bootstrap.servers=${KAFKA_CLUSTER}-kafka-bootstrap:9092 --record-size=10 --throughput 1 --num-records 1000 --topic ${TOPIC_NAME}
```

Kafka Consumer and Consumer Group

Console consumer (without a consumer group)

```
# using a new pod
oc run kafka-consumer -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true --restart=Never -- bin/kafka-console-consumer.sh --bootstrap-server ${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --topic ${TOPIC_NAME}

# using kafka-0 broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -- env - bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --from-beginning --topic ${TOPIC_NAME}
```

Use the “--whitelist” option to subscribe to multiple topics using regex instead of --topic option, for example --whitelist topicprefix.*

Console consumer (with a consumer group)

```
# using a new pod
oc run kafka-consumer -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true --restart=Never -- bin/kafka-console-consumer.sh --bootstrap-server ${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --topic ${TOPIC_NAME} --group ${CONSUMER_GROUP_NAME}
```

```
# using kafka-0 broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -- bin/kafka-console-consumer.sh --bootstrap-server
localhost:9092 --from-beginning --topic ${TOPIC_NAME} --group ${CONSUMER_GROUP_NAME}
```

Consumer performance test tool

```
oc run kafka-consumer-perf -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
--rm=true --restart=Never -- bin/kafka-consumer-perf-test.sh --broker-list
${KAFKA_CLUSTER}-kafka-bootstrap:9092 --threads 1 --from-latest --print-metrics
--show-detailed-stats=true --reporting-interval 1000 --timeout 100000000 --messages 10000 --topic
${TOPIC_NAME}
```

Consumer group list

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --list
```

Consumer group description (--offset)

```
# a specific consumer group
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --describe --group ${CONSUMER_GROUP_NAME}
~~~
GROUP          TOPIC          PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG          CONSUMER-ID
HOST          CLIENT-ID
my-group      test_topic      0           30              30              0
consumer-1-0f11a1f8-8f0c-4cd3-9499-471b62332d7f /10.128.0.202 consumer-1
my-group      test_topic      4           30              30              0
consumer-1-0f11a1f8-8f0c-4cd3-9499-471b62332d7f /10.128.0.202 consumer-1
~~~

# all the consumer groups
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --describe --all-groups
```

Consumer group description (--state)

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --describe --group ${CONSUMER_GROUP_NAME} --state
~~~
GROUP          COORDINATOR (ID)
ASSIGNMENT-STRATEGY  STATE          #MEMBERS
my-group
my-cluster-kafka-2.my-cluster-kafka-brokers.my-project.svc:9092 (2) range      Stable
1
~~~
```

Delete a consumer group

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --delete --group ${CONSUMER_GROUP_NAME}
```

Reset the offsets for a specific consumer group

```
# rest to earliest
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --reset-offsets --to-earliest --group ${CONSUMER_GROUP_NAME}
--topic ${TOPIC_NAME} --execute
```

```
# rest to specific offset
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --reset-offsets --to-offset <OFFSET_VALUE> --group
${CONSUMER_GROUP_NAME} --topic ${TOPIC_NAME}:<PARTITION_NUMBERS> --execute
```

Don't forget `--execute` option. For the other options on how to reset other than earliest, see [here](#).

Delete offset with specific topic for a specific consumer group

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh
--bootstrap-server localhost:9092 --delete-offsets --group ${CONSUMER_GROUP_NAME} --topic
${TOPIC_NAME}
```

About Kafka CLI execution environment

Executing Kafka CLI using an existing broker

```
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - <your_command>
```

This is the easiest way to run the Kafka CLI such as `bin/kafka-topics.sh`. For this reason, It is also used as example commands above. However, If you don't want to burden the broker, you should create a new pod to execute commands.

[IMPORTANT] when running the Kafka CLI on a broker, you should add "env -" to remove the environment variable because even "-Xms" specified for the Kafka broker could be applied for the CLI commands and it can lead out of memory error.

For "oc exe" using the existing kafka broker pods, you can use localhost:9092 as a Kafka broker, or localhost:2181 to Zookeeper via `tls-sidecar` (until AMQ Streams 1.5) as below:

```
# access Kafka
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --bootstrap-server
localhost:9092 --list

# access Zookeeper
# since AMQ Stream 1.6
oc exec -it ${KAFKA_CLUSTER}-zookeeper-0 -- env - bin/zookeeper-shell.sh 127.0.0.1:2181 delete
/controller
# via local tls-sidecar: until AMQ Streams 1.5
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/zookeeper-shell.sh zoo-entrance:2181 ls
/brokers/ids
```

Executing Kafka CLI using a new pod (do not use a pod of an existing broker)

```
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- <your_command>
```

This method does not allow to use localhost:9092 as a Kafka broker, or localhost:2181 to Zookeeper via `tls-sidecar`. You can use service to access kafka, and use [zoo-entrance](#) to access Zookeeper, and so on like below:

```
# access Kafka via bootstrap service
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- bin/kafka-topics.sh --bootstrap-server ${KAFKA_CLUSTER}-kafka-bootstrap:9092
--describe ${TOPIC_NAME}

# FYI, access Zookeeper via zoo-entrance
# until AMQ Stream 1.5
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- bin/zookeeper-shell.sh zoo-entrance:2181 ls /brokers/ids
```

Executing Kafka CLI using a customized new pod as root

```
# create imagestreams of amq-streams image
oc import-image amq-streams-debug --from=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
--confirm

# run a new pod based on amq-streams image as root and execute your commands
oc debug istag/amq-streams-debug:latest --as-root=true
> <your command>

# FYI, you can delete the imagestream for cleanup after debugging
oc delete imagestream amq-streams-debug
```

It can be used to execute kafka CLI such as `bin/kafka-topics.sh` in any namespace. You can also install various packages using `microdnf`.

Additionally, you can customize the amq-streams Kafka image using Dockerfile as below:

```
# step 1: create Dockerfile that extends amq-streams image.
vi Dockerfile
~~~
FROM registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
USER root:root
RUN microdnf install -y curl java-1.8.0-openjdk-devel
USER 1001
~~~

# step 2: Docker build on the OpenShift
oc new-build --name=amq-streams-debug --strategy=docker --binary
oc start-build amq-streams-debug --from-dir=. --follow

# step 3: run the image as root and execute your commands
oc debug istag/amq-streams-debug:latest --as-root=true
> <your command>

# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l app=amq-streams-debug -l build=amq-streams-debug
```

Strimzi Cluster Operator

Strimzi Cluster Operator definition

```
oc get deployment/strimzi-cluster-operator -o yaml

# internal Cluster Operator configuration
oc exec deployment/strimzi-cluster-operator -- ps aux
```

Stop cluster operator

```
# stop the cluster operator
oc scale deployment/strimzi-cluster-operator --replicas=0

# restart the cluster operator
oc scale deployment/strimzi-cluster-operator --replicas=1
```

Pause reconciliation of custom resources

```
# pause reconciliation since AMQ Streams 1.7, for topic resource since AMQ Streams 1.8
oc annotate KIND-OF-CUSTOM-RESOURCE NAME-OF-CUSTOM-RESOURCE strimzi.io/pause-reconciliation="true"
```

```
# restart reconciliation
oc annotate KIND-OF-CUSTOM-RESOURCE NAME-OF-CUSTOM-RESOURCE strimzi.io/pause-reconciliation-
```

Check for duplicate cluster operators installed

```
oc get pods --all-namespaces
oc get pods --all-namespaces -l "strimzi.io/kind=cluster-operator" -o yaml

oc get packagemanifests -n openshift-marketplace
oc get packagemanifests/amq-streams -n openshift-marketplace -o yaml
oc get packagemanifests/strimzi-kafka-operator -n openshift-marketplace -o yaml

oc get Subscription --all-namespaces -o yaml
```

Strimzi Entity Operator (comprises Topic Operator and User Operator)

Strimzi Entity Operator definition

```
oc get Kafka -o yaml
oc get deployment/${KAFKA_CLUSTER}-entity-operator -o yaml

# internal Entity Operator configuration
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- ps aux
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- ps aux
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- cat
/opt/topic-operator/custom-config/log4j2.properties
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- cat
/opt/user-operator/custom-config/log4j2.properties
```

Stop entity operator

```
# To stop entity operator, comment out "entityOperator" and apply it.
~~~
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  kafka:
    # entityOperator:
    #   topicOperator: {}
    #   userOperator: {}
~~~

# FYI, you CANNOT stop the entity operator using the following command because CO will revert it.
oc scale deployment/${KAFKA_CLUSTER}-entity-operator --replicas=0
```

Kafka Connect and Debezium

Create a MySQL DB POD for test

```
# create a mysql pod
oc new-app --name=mysql debezium/example-mysql:1.0 -e MYSQL_ROOT_PASSWORD=debezium -e
MYSQL_USER=mysqluser -e MYSQL_PASSWORD=mysqlpw -l name=mysql-debug

# login to mysql cli
oc exec deployment/mysql -it -- mysql -umysqluser -pmysqlpw inventory
```



```
# execute mysql commands you want
show tables;
select * from customers;
insert into customers values (default, "Sarah", "Thompson", "kitt@acme.com");

# FYI, for cleanup after debugging
oc delete all -l name=mysql-debug
```

Create Kafka Connect image with Debezium: build a image on the OpenShift

```
# step 1
mkdir ./my-plugins

# step 2: download kafka connect plugins or Debezium plugins
https://access.redhat.com/jbossnetwork/restricted/listSoftware.html?product=red.hat.integration&downloadType=distributions

# step 3: extract files and put it in ./my-plugins

$ tree ./my-plugins/
./my-plugins/
├── debezium-connector-db2
│   ├── ...
├── debezium-connector-mongodb
│   ├── ...
├── debezium-connector-mysql
│   ├── ...
├── debezium-connector-postgres
│   ├── ...
├── debezium-connector-sqlserver
│   ├── ...
```

```
# step 4: create Dockerfile that extends amq-streams image.
vi Dockerfile
~~~
FROM registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
USER root:root
COPY ./my-plugins/ /opt/kafka/plugins/
USER 1001
~~~

# step 5: Docker build on the OpenShift
oc new-build --name=my-kafka-connect --strategy=docker --binary
oc start-build my-kafka-connect --from-dir=. --follow

# step 6: set KafkaConnector.spec.image to the new image path and apply.
vi examples/kafka-connect/kafka-connect-single-node-kafka.yaml
~~~
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnector
metadata:
  name: my-connect-cluster
spec:
  #...
  image: image-registry.openshift-image-registry.svc:5000/<$$namespace$$>/my-kafka-connect:latest
~~~

# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l build=my-kafka-connect
```

Create a MySQL Debezium connector accessing to MySQL DB on the Kafka Connector POD

```
oc exec deployment/${KAFKACONNECT_CLUSTER}-connect -- curl -s -v -X POST -H
"Accept:application/json" -H "Content-Type:application/json" localhost:8083/connectors/ -d '{
"name": "inventory-connector", "config": { "connector.class":
"io.debezium.connector.mysql.MySqlConnector", "tasks.max": "1", "database.hostname": "mysql",
"database.port": "3306", "database.user": "debezium", "database.password": "dbz",
"database.server.id": "184054", "database.server.name": "dbserver1", "database.whitelist":
"inventory", "database.history.kafka.bootstrap.servers": "${KAFKA_CLUSTER}'-kafka-bootstrap:9092",
"database.history.kafka.topic": "schema-changes.inventory" } }'

# FYI, you can subscribe "dbserver1.inventory.*" topics using the following command.
oc run kafka-consumer -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- bin/kafka-console-consumer.sh --bootstrap-server
${KAFKA_CLUSTER}-kafka-bootstrap:9092 --from-beginning --whitelist 'dbserver1.inventory.*'

# FYI, you can insert a row to MySQL DB using the following command.
oc exec deployment/mysql -it -- mysql -umysqluser -pmysqlpw inventory -e 'insert into customers
values (default, "Sarah", "Thompson", "kitt@acme.com")'
```

Cruise Control

Cruise Control definition

```
oc get Kafka -o yaml
oc get deployment/${KAFKA_CLUSTER}-cruise-control -o yaml

# internal Kafka configuration
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- cat /tmp/cruisecontrol.properties
# log4j configuration since AMQ Streams 1.7
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- cat
/opt/cruise-control/custom-config/log4j.properties
# log4j configuration until AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- cat
/opt/cruise-control/custom-config/log4j.properties
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- ps aux
```

TLS

Validity days of Cluster CA and Clients CA certifications in secrets

```
# Cluster CA
oc get secret ${KAFKA_CLUSTER}-cluster-ca-cert -o 'jsonpath={.data.ca\.crt}' | base64 -d | openssl
x509 -subject -issuer -startdate -enddate -noout
~~~
subject=O = io.strimzi, CN = cluster-ca v0
issuer=O = io.strimzi, CN = cluster-ca v0
notBefore=Jun 30 09:43:54 2020 GMT
notAfter=Jun 30 09:43:54 2021 GMT
~~~

# Clients CA
oc get secret ${KAFKA_CLUSTER}-clients-ca-cert -o 'jsonpath={.data.ca\.crt}' | base64 -d | openssl
x509 -subject -issuer -startdate -enddate -noout
~~~
```

```

subject=0 = io.strimzi, CN = clients-ca v0
issuer=0 = io.strimzi, CN = clients-ca v0
notBefore=Jun 30 09:43:54 2020 GMT
notAfter=Jun 30 09:43:54 2021 GMT
~~~

# FYI:
# Kafka Broker
oc get secret ${KAFKA_CLUSTER}-kafka-brokers -o 'jsonpath={.data.'${KAFKA_CLUSTER}'-kafka-0\.crt}' |
base64 -d | openssl x509 -subject -issuer -startdate -enddate -noout
# Zookeeper
oc get secret ${KAFKA_CLUSTER}-zookeeper-nodes -o
'jsonpath={.data.'${KAFKA_CLUSTER}'-zookeeper-0\.crt}' | base64 -d | openssl x509 -subject -issuer
-startdate -enddate -noout
# Entity Operator
oc get secret ${KAFKA_CLUSTER}-entity-operator-certs -o 'jsonpath={.data.entity-operator\.crt}' |
base64 -d | openssl x509 -subject -issuer -startdate -enddate -noout

```

Check actual TLS info using openssl s_client command

```

# check to connect to server
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- openssl s_client -debug -connect
<TLS_URL_ADDRESS>:<TLS_PORT_NUMBER> -servername <TLS_URL_ADDRESS>

# show validity days
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- openssl s_client -debug -connect
<TLS_URL_ADDRESS>:<TLS_PORT_NUMBER> -servername <TLS_URL_ADDRESS> 2>/dev/null | openssl x509
-subject -issuer -startdate -enddate -noout

```

Check TLS port on a Kafka broker using openssl s_client command

```

oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- openssl s_client -debug -connect
localhost:<TLS_PORT_NUMBER> < /dev/null

```

Enable ssl debug log on Kafka brokers

```

~~~
Kind: Kafka
spec:
  kafka:
    template:
      kafkaContainer:
        env:
          - name: KAFKA_DEBUG
            value: "true"
          - name: JAVA_DEBUG_OPTS
            value: "-Djavax.net.debug=ssl" # or "-Djavax.net.debug=all"
~~~

```

Renewing CA certificates manually

```

# cluster CA
oc annotate secret ${KAFKA_CLUSTER}-cluster-ca-cert strimzi.io/force-renew=true

# clients CA
oc annotate secret ${KAFKA_CLUSTER}-clients-ca-cert strimzi.io/force-renew=true

```

for details, refer to the [document](#).

Replace to my own cluster ca certification and its key

```
# I recommend recreating the kafka cluster before replacing to my own cluster ca.
oc delete kafka/${KAFKA_CLUSTER}

# set Kafka.clusterCa.generateCertificateAuthority: false
# vi kafka.yaml
~~~
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
...
  clusterCa:
    generateCertificateAuthority: false
~~~

# replace cluster ca
mkdir cert
cd cert

openssl genrsa 2024 > ca.key
openssl req -new -key ca.key -subj
"/C=JP/ST=Tokyo/L=Shibuya-ku/O=redhatexample/OU=cee/CN=amqstreams.redhatexample.com" > ca.csr
openssl x509 -req -days 365 -signkey ca.key < ca.csr > ca.crt

oc delete secret ${KAFKA_CLUSTER}-cluster-ca-cert
oc create secret generic ${KAFKA_CLUSTER}-cluster-ca-cert --from-file=ca.crt=./ca.crt

oc delete secret ${KAFKA_CLUSTER}-cluster-ca
oc create secret generic ${KAFKA_CLUSTER}-cluster-ca --from-file=ca.key=./ca.key

oc label secret ${KAFKA_CLUSTER}-cluster-ca-cert strimzi.io/kind=Kafka strimzi.io/cluster=${KAFKA_CLUSTER}
oc label secret ${KAFKA_CLUSTER}-cluster-ca strimzi.io/kind=Kafka strimzi.io/cluster=${KAFKA_CLUSTER}

# for cleanup
cd ..
rm -R ./cert

# apply kafka.yaml
oc apply -f kafka.yaml
```

If you want to test cluster CA certification expiration, you can set “-days 0”.

Enable SSL debug log of Kafka CLI for debugging

```
# Kafka console producer
KAFKA_OPTS=-Djavax.net.debug=ssl bin/kafka-console-producer.sh ...

# Kafka producer performance test tool
KAFKA_OPTS=-Djavax.net.debug=ssl bin/kafka-producer-perf-test.sh ...

# Kafka console consumer
KAFKA_OPTS=-Djavax.net.debug=ssl bin/kafka-console-consumer.sh ...
```

Use KAFKA_OPTS=-Djavax.net.debug=ssl

Cluster CA and truststore.p12

```
# Cluster CA
oc get secret ${KAFKA_CLUSTER}-cluster-ca-cert -o jsonpath='{.data.ca\.crt}' | base64 -d > ca.crt

# truststore.p12
keytool -keystore truststore.p12 -storepass <your truststore password> -noprompt -alias ca -import
-file ca.crt -storetype PKCS12
```

KafkaUser secret and keystore.p12

```
# KafkaUser
oc get secret ${KAFKAUSER_NAME} -o jsonpath='{.data.user\.p12}' | base64 -d > user.p12
oc get secret ${KAFKAUSER_NAME} -o jsonpath='{.data.user\.password}' | base64 -d > user.password

# keystore.p12
keytool -importkeystore -destkeystore keystore.p12 -srckeystore user.p12 -srcstorepass $(cat
user.password) -srcstoretype PKCS12 -deststoretype PKCS12 -destkeypass <your keystore password>
-deststorepass <your keystore password as same as "destkeypass">;
```

tls listener (port:9093) with one way TLS

Prerequisites

```
# modify kafka.yaml and apply it
~~~
#   since AMQ Streams 1.6
kind: Kafka
spec:
  kafka:
    listeners:
      ...
      - name: tls
        port: 9093
        type: internal
        tls: true

#   until AMQ Streams 1.5
kind: Kafka
spec:
  kafka:
    listeners:
      tls: {}
      ...
~~~
```

Console producer

```
oc run kafka-producer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{"name": "cluster-ca", "secret": {"secretName":
"${KAFKA_CLUSTER}-cluster-ca-cert"}}, {"name": "workdir", "emptyDir": {}}], "containers": [{
"name": "kafka-producer-tls", "image": "registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty": true, "volumeMounts": [{"mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{"mountPath": "/opt/kafka/workdir", "name": "workdir"}], "command": ["bash", "-c", "keytool -keystore
workdir/truststore.p12 -storepass password -noprompt -alias ca -import -file
/opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12; bin/kafka-console-producer.sh --broker-list
`${KAFKA_CLUSTER}`-kafka-bootstrap:9093 --producer-property security.protocol=SSL
--producer-property ssl.truststore.location=workdir/truststore.p12 --producer-property
ssl.truststore.password=password --producer-property ssl.truststore.type=PKCS12 --topic
`${TOPIC_NAME}` " ] } ] }'
```

Producer performance test tool (size:10, throughput:1, records:1000)

```
oc run kafka-producer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{"name": "cluster-ca", "secret": {"secretName":
"${KAFKA_CLUSTER}-cluster-ca-cert"}}, {"name": "workdir", "emptyDir": {}}], "containers": [{
```

```
"name": "kafka-producer-tls", "image": "registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty": true, "volumeMounts": [{ "mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{ "mountPath": "/opt/kafka/workdir", "name": "workdir"}], "command": ["bash", "-c", "keytool -keystore
workdir/truststore.p12 -storepass password -noprompt -alias ca -import -file
/opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12; bin/kafka-producer-perf-test.sh
--producer-props bootstrap.servers='${KAFKA_CLUSTER}'-kafka-bootstrap:9093 security.protocol=SSL
ssl.truststore.location=workdir/truststore.p12 ssl.truststore.password=password
ssl.truststore.type=PKCS12 --record-size=10 --throughput 1 --num-records 1000 --topic
'${TOPIC_NAME}' " ] ] ] }
```

Console consumer

```
oc run kafka-consumer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{ "name": "cluster-ca", "secret": { "secretName":
"${KAFKA_CLUSTER}-cluster-ca-cert"} }, { "name": "workdir", "emptyDir": {} } ], "containers": [{
"name": "kafka-consumer-tls", "image": "registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty": true, "volumeMounts": [{ "mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{ "mountPath": "/opt/kafka/workdir", "name": "workdir"}], "command": ["bash", "-c", "keytool -keystore
workdir/truststore.p12 -storepass password -noprompt -alias ca -import -file
/opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12; bin/kafka-console-consumer.sh
--bootstrap-server '${KAFKA_CLUSTER}'-kafka-bootstrap:9093 --consumer-property security.protocol=SSL
--consumer-property ssl.truststore.location=workdir/truststore.p12 --consumer-property
ssl.truststore.password=password --consumer-property ssl.truststore.type=PKCS12 --from-beginning
--topic '${TOPIC_NAME}' " ] ] ] }
```

tls listener (port:9093) with two-way(mutual) TLS client authentication and simple authorization

Prerequisites

```
# modify kafka.yaml and apply it
~~~
# since AMQ Streams 1.6
kind: Kafka
spec:
  kafka:
    listeners:
      ...
      - name: tls
        port: 9093
        type: internal
        tls: true
        authentication:
          type: tls
      ...
    authorization:
      type: simple

# until AMQ Streams 1.5
kind: Kafka
spec:
  kafka:
    listeners:
      tls:
        authentication:
          type: tls
    authorization:
```

```

    type: simple
    ...
~~~
# modify examples/user/kafka-user.yaml and apply it
~~~
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: my-user
...
spec:
  authentication:
    type: tls
...

```

Console producer

```

oc run kafka-producer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{"name": "cluster-ca", "secret": {"secretName":
"${KAFKA_CLUSTER}-cluster-ca-cert"}}, {"name": "kafkauser", "secret": {"secretName":
"${KAFKAUSER_NAME}"}}], {"name": "workdir", "emptyDir": {}}, {"containers": [{ "name":
"kafka-producer-tls", "image": "registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty": true, "volumeMounts": [{"mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{"mountPath": "/opt/kafka/kafkauser", "name": "kafkauser"}, {"mountPath": "/opt/kafka/workdir",
"name": "workdir"}], "command": ["bash", "-c", "keytool -keystore workdir/truststore.p12 -storepass
password -noprompt -alias ca -import -file /opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12;
keytool -importkeystore -destkeystore workdir/keystore.p12 -srckeystore
/opt/kafka/kafkauser/user.p12 -srcstorepass $(cat /opt/kafka/kafkauser/user.password) -srcstoretype
PKCS12 -deststoretype PKCS12 -destkeypass password -deststorepass password;
bin/kafka-console-producer.sh --broker-list '${KAFKA_CLUSTER}'-kafka-bootstrap:9093
--producer-property security.protocol=SSL --producer-property
ssl.truststore.location=workdir/truststore.p12 --producer-property ssl.truststore.password=password
--producer-property ssl.truststore.type=PKCS12 --producer-property
ssl.keystore.location=workdir/keystore.p12 --producer-property ssl.keystore.password=password
--producer-property ssl.keystore.type=PKCS12 --topic '${TOPIC_NAME}' "]" ]} ]}'

```

Producer performance test tool (size:10, throughput:1, records:1000)

```

oc run kafka-producer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{"name": "cluster-ca", "secret": {"secretName":
"${KAFKA_CLUSTER}-cluster-ca-cert"}}, {"name": "kafkauser", "secret": {"secretName":
"${KAFKAUSER_NAME}"}}], {"name": "workdir", "emptyDir": {}}, {"containers": [{ "name":
"kafka-producer-tls", "image": "registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty": true, "volumeMounts": [{"mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{"mountPath": "/opt/kafka/kafkauser", "name": "kafkauser"}, {"mountPath": "/opt/kafka/workdir",
"name": "workdir"}], "command": ["bash", "-c", "keytool -keystore workdir/truststore.p12 -storepass
password -noprompt -alias ca -import -file /opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12;
keytool -importkeystore -destkeystore workdir/keystore.p12 -srckeystore
/opt/kafka/kafkauser/user.p12 -srcstorepass $(cat /opt/kafka/kafkauser/user.password) -srcstoretype
PKCS12 -deststoretype PKCS12 -destkeypass password -deststorepass password;
bin/kafka-producer-perf-test.sh --producer-props
bootstrap.servers='${KAFKA_CLUSTER}'-kafka-bootstrap:9093 security.protocol=SSL
ssl.truststore.location=workdir/truststore.p12 ssl.truststore.password=password
ssl.truststore.type=PKCS12 ssl.keystore.location=workdir/keystore.p12 ssl.keystore.password=password
ssl.keystore.type=PKCS12 --record-size=10 --throughput 1 --num-records 1000 --topic '${TOPIC_NAME}'
"] ]} ]}'

```

Console consumer (with a consumer group)

```
oc run kafka-consumer-tls -ti --image=ignore --rm=true --restart=Never --overrides='{ "apiVersion":
"v1", "spec": { "volumes": [{"name": "cluster-ca","secret": {"secretName":
"${KAFKA_CLUSTER}"-cluster-ca-cert"}}, {"name": "kafkauser","secret": {"secretName":
"${KAFKAUSER_NAME}"}}], {"name": "workdir","emptyDir": {}}, "containers": [{ "name":
"kafka-consumer-tls", "image":"registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0",
"tty":true, "volumeMounts": [{"mountPath": "/opt/kafka/cluster-ca-certs", "name": "cluster-ca"},
{"mountPath": "/opt/kafka/kafkauser", "name": "kafkauser"}, {"mountPath": "/opt/kafka/workdir",
"name": "workdir"}], "command":["bash","-c", "keytool -keystore workdir/truststore.p12 -storepass
password -noprompt -alias ca -import -file /opt/kafka/cluster-ca-certs/ca.crt -storetype PKCS12;
keytool -importkeystore -destkeystore workdir/keystore.p12 -srckeystore
/opt/kafka/kafkauser/user.p12 -srcstorepass $(cat /opt/kafka/kafkauser/user.password) -srcstoretype
PKCS12 -deststoretype PKCS12 -destkeypass password -deststorepass password;
bin/kafka-console-consumer.sh --bootstrap-server '${KAFKA_CLUSTER}'-kafka-bootstrap:9093
--consumer-property security.protocol=SSL --consumer-property
ssl.truststore.location=workdir/truststore.p12 --consumer-property ssl.truststore.password=password
--consumer-property ssl.truststore.type=PKCS12 --consumer-property
ssl.keystore.location=workdir/keystore.p12 --consumer-property ssl.keystore.password=password
--consumer-property ssl.keystore.type=PKCS12 --from-beginning --topic '${TOPIC_NAME}' --group
'${CONSUMER_GROUP_NAME}' " ] } ] }'
```

JMX

Kafka provides many metrics via JMX. For details of mbeans Kafka has for metrics, refer to the [document](#).

Enable JMX port(9999) on the Kafka brokers

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    ...
    jmxOptions: {}
```

For details of jmxOptions, refer to the [document](#). jmxOptions is available since AMQ Streams 1.4.

Print Kafka metrics via JMX using kafka.tools.JmxTool

```
# using an existing broker
oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-run-class.sh kafka.tools.JmxTool
--reporting-interval 1000 --object-name kafka.server:type=BrokerTopicMetrics,name=MessagesInPerSec

# using a new pod (connect via headless service)
oc run kafka-exe -ti --image=registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0 --rm=true
--restart=Never -- bin/kafka-run-class.sh kafka.tools.JmxTool --jmx-url
service:jmx:rmi:///jndi/rmi://${KAFKA_CLUSTER}-kafka-${BROKER_ID}.${KAFKA_CLUSTER}-kafka-brokers:999
9/jmxrmi --reporting-interval 1000 --object-name
kafka.server:type=BrokerTopicMetrics,name=MessagesInPerSec
```

This method can display metrics through JMX. To output only one time you can use the "--one-time true" option.

jmxterm

```
# get jmxterm tool jar
curl -L -O https://github.com/jiaqi/jmxterm/releases/download/v1.0.2/jmxterm-1.0.2-uber.jar
```



```
# run interactive jmxterm
java -jar jmxterm-1.0.2-uber.jar -l <jmx address>:<jmx port>
```

Jmxterm is an open source command line based interactive JMX client written in Java. It lets user access a Java MBean server in command line console without graphical environment. In another word, it's a command line alternative of [jconsole](#). For details of jmxterm, refer to [jmxterm web page](#).

jconsole (or jmxterm) from localhost machine to a Kafka broker

```
# add loopback address of the same IP address as a Kafka broker on the localhost machine
sudo ip addr add $(oc get pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} -o=jsonpath='{.status.podIP}')/32
dev lo

# open port forwarding to a Kafka broker on the localhost machine
oc port-forward pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} 9999:9999 --address=$(oc get
pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} -o=jsonpath='{.status.podIP}')

# [on another terminal] execute jconsole
jconsole $(oc get pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} -o=jsonpath='{.status.podIP}'):9999
# Or, [on another terminal] execute jmxterm
java -jar jmxterm-1.0.2-uber.jar -l $(oc get pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID}
-o=jsonpath='{.status.podIP}'):9999

# FYI, delete loopback address you've created for cleanup
sudo ip addr del $(oc get pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} -o=jsonpath='{.status.podIP}')/32
dev lo
```

This has the advantage of not polluting the AMQ Streams environment, but you need to create a loopback address on the localhost machine. " -Djava.rmi.server.hostname=XX.XX.XX.XX" is specified in Kafka broker so If you don't use that loopback IP address to access the Kafka broker, the connection will be denied.

jmxterm from a new pod to a Kafka broker

```
oc run jmxterm -ti --image=registry.redhat.io/rhel8/support-tools -- bash -c "dnf install -y curl
java-1.8.0-openjdk; curl -L -O
https://github.com/jiaqi/jmxterm/releases/download/v1.0.2/jmxterm-1.0.2-uber.jar; java -jar
jmxterm-1.0.2-uber.jar -l ${KAFKA_CLUSTER}-kafka-${BROKER_ID}.${KAFKA_CLUSTER}-kafka-brokers:9999"
```

This method is easy, but it takes time because it downloads the java package.

JMX Exporter and Kafka Exporter

To enable JMX Exporter, for details please refer to [the AMQ Streams document](#). You can also use a sample Kafka yaml (please adjust the version of the YAML file in the GitHub).

Component	Custom resource	Example YAML file
Kafka and ZooKeeper	Kafka	kafka-metrics.yaml
Kafka Connect	KafkaConnect and KafkaConnectS2I	kafka-connect-metrics.yaml
Kafka MirrorMaker 2.0	KafkaMirrorMaker2	kafka-mirror-maker-2-metrics.yaml
Kafka Bridge	KafkaBridge	kafka-bridge-metrics.yaml
Cruise Control	Kafka	kafka-cruise-control-metrics.yaml

Show JMX exporter metrics using curl

```
# Kafka broker
oc exec -it ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- curl http://localhost:9404/metrics

# Zookeeper
oc exec -it ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -- curl http://localhost:9404/metrics
```

Show Kafka exporter metrics using curl

```
# since AMQ Streams 1.6
oc exec -it deployment/${KAFKA_CLUSTER}-kafka-exporter -- curl http://localhost:9404/metrics
```

Log

Log of a pod

```
# Kafka broker
oc logs ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka
oc logs -p ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka
oc logs ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c tls-sidecar
oc logs -p ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c tls-sidecar

# Zookeeper
oc logs ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper
oc logs -p ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper
oc logs ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c tls-sidecar
oc logs -p ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c tls-sidecar

# Cluster operator
oc logs deployment/strimzi-cluster-operator
oc logs -p deployment/strimzi-cluster-operator

# Entity operator
oc logs deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator
oc logs -p deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator
oc logs deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator
oc logs -p deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator
oc logs deployment/${KAFKA_CLUSTER}-entity-operator -c tls-sidecar
oc logs -p deployment/${KAFKA_CLUSTER}-entity-operator -c tls-sidecar

# Cruise Control
oc logs deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control
oc logs -p deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control
oc logs deployment/${KAFKA_CLUSTER}-cruise-control -c tls-sidecar
oc logs -p deployment/${KAFKA_CLUSTER}-cruise-control -c tls-sidecar

# mirror maker : It only works when there is only one pod because it grep the pod name.
oc logs $(oc get pods -l strimzi.io/cluster -o name | grep mirror-maker-)
oc logs -p $(oc get pods -l strimzi.io/cluster -o name | grep mirror-maker-)

# Kafka exporter : It only works when there is only one pod because it grep the pod name.
oc logs $(oc get pods -l strimzi.io/cluster -o name | grep kafka-exporter-)
oc logs -p $(oc get pods -l strimzi.io/cluster -o name | grep kafka-exporter-)
```

```
# Kafka connect: It only works when there is only one pod because it grep the pod name.
oc logs $(oc get pods -l strimzi.io/cluster -o name | grep connect-)
oc logs -p $(oc get pods -l strimzi.io/cluster -o name | grep connect-)

# Kafka bridge: It only works when there is only one pod because it grep the pod name.
oc logs $(oc get pods -l strimzi.io/cluster -o name | grep bridge-)
oc logs -p $(oc get pods -l strimzi.io/cluster -o name | grep bridge-)
```

Use [dump script](#) for a comprehensive log output.

Use the "-f" option to keep to follow log.

Use the "-p" option to get the previous log.

Multiple pod logs using stern

```
# Kafka brokers and Zookeepers logs
stern "${KAFKA_CLUSTER}-(kafka|zookeeper)" --tail 0

# cluster operator and entity-operator logs
stern "(strimzi-cluster-operator|entity-operator)" --tail 0

# cluster operator, entity-operator, Kafka brokers, Zookeepers logs
stern "(strimzi-cluster-operator|entity-operator|${KAFKA_CLUSTER}-(kafka|zookeeper))" --tail 0
```

To install stern, refer to <https://github.com/wercker/stern>.

Enable debug log : Kafka, Zookeeper, Entity Operator, Cruise Control and its tls-sidecar

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
spec:
  # ...
  kafka:
    # ...
    logging:
      type: inline
      loggers:
        kafka.root.logger.level: "DEBUG"
        log4j.logger.kafka: "DEBUG"
        log4j.logger.org.apache.kafka: "DEBUG"
    tlsSidecar:
      # ...
      logLevel: debug
  # ...
  zookeeper:
    # ...
    logging:
      type: inline
      loggers:
        zookeeper.root.logger: "DEBUG"
    tlsSidecar:
      # ...
      logLevel: debug
  # ...
  entityOperator:
    # ...
    topicOperator:
      # ...
      logging:
        type: inline
        loggers:
```

```

        rootLogger.level: DEBUG
# ...
userOperator:
# ...
  logging:
    type: inline
    loggers:
      rootLogger.level: DEBUG
# ...
tlsSidecar:
# ...
  logLevel: debug
# ...
cruiseControl:
# ...
  logging:
    type: inline
    loggers:
      # since AMQ Streams 1.7
      rootLogger.level: "DEBUG"
      # until AMQ Streams 1.6
      #   you need to restart Cruise Control pod after changing about 5 minute later.
      #   For details, please refer to ENTMQST-2654
      cruisecontrol.root.logger: "DEBUG"
tlsSidecar:
# ...
  logLevel: debug

# FYI, you can confirm actual log4j file
# actual Kafka log4j file
oc exec ${KAFKA_CLUSTER}-kafka-0 -- cat /opt/kafka/custom-config/log4j.properties

# actual Zookeeper log4j file
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- cat /opt/kafka/custom-config/log4j.properties

# actual topic operator log4j file in entity operator
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- cat
/opt/topic-operator/custom-config/log4j2.properties

# actual user operator log4j file in entity operator
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- cat /opt/usecd
plr-operator/custom-config/log4j2.properties

# actual cruise control log4j file
#   until AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- cat
/opt/cruise-control/custom-config/log4j.properties
#   since AMQ Streams 1.7
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -- cat
/opt/cruise-control/custom-config/log4j2.properties

```

The above is an example of setting the log level to DEBUG for root loggers. Enabling the DEBUG log levels results in a large amount of log output. Change the log level only where necessary. Note that several loggers may not be affected by the above settings because they have a log level specified by default. For details, refer to the [document](#) or actual log4j file.

Enable debug log : Cluster Operator

```

# since AMQ Streams 1.6
## pattern 1 : change STRIMZI_LOG_LEVEL environment variable value
vi install/cluster-operator/060-Deployment-strimzi-cluster-operator.yaml
~~~

```

```

spec:
  ...
  template:
    ...
    spec:
      ...
      containers:
        ...
        env:
          ...
          - name: STRIMZI_LOG_LEVEL
            value: "DEBUG"
    ~~~

## pattern 2 : change log configuration on the config map
vi install/cluster-operator/050-ConfigMap-strimzi-cluster-operator.yaml
~~~
data:
  log4j2.properties: |
    ...
    #rootLogger.level = ${env:STRIMZI_LOG_LEVEL:-INFO} # comment out
    rootLogger.level = DEBUG
  ~~~

# until AMQ Streams 1.5
vi install/cluster-operator/050-Deployment-strimzi-cluster-operator.yaml
~~~
spec:
  template:
    spec:
      containers:
        - name: strimzi-cluster-operator
          env:
            - name: STRIMZI_LOG_LEVEL
              value: "DEBUG"
    ~~~

```

Enable debug log : Kafka Bridge

```

apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaBridge
spec:
  # ...
  logging:
    type: inline
    loggers:
      bridge.root.logger: "DEBUG"
      log4j.logger.io.strimzi.kafka.bridge: "DEBUG"

# FYI, you can confirm actual log4j file
oc exec $(oc get pods -l strimzi.io/cluster -o name | grep bridge- | head -n 1) -- cat
/opt/strimzi/custom-config/log4j.properties

```

The default configuration for healthy and ready endpoints is WARN level.

Enable debug log : Kafka Connect

```

# when to use CR
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaConnect
spec:

```

```

# ...
logging:
  type: inline
  loggers:
    connect.root.logger.level: "DEBUG"
# ...

# FYI, you can confirm actual log4j configuration
oc exec deployment/${KAFKACONNECT_CLUSTER}-connect -- curl http://localhost:8083/admin/loggers/
# Or
oc exec deployment/${KAFKACONNECT_CLUSTER}-connect -- cat /opt/kafka/custom-config/log4j.properties

```

Same for "Kind: KafkaConnectS2I".

Enable debug log : Kafka Mirror Maker v1

```

apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker
spec:
  # ...
  logging:
    type: inline
    loggers:
      mirrormaker.root.logger: "DEBUG"
  # ...

```

Enable debug log : Kafka Mirror Maker v2

```

apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaMirrorMaker2
spec:
  # ...
  logging:
    type: inline
    loggers:
      connect.root.logger.level: "DEBUG"
  # ...

```

Kafka Client Example

For details, refer to [Apache Kafka client examples](#).

Run Kafka producer client example on the OpenShift

```
# step 1: clone Kafka Client Example
git clone https://github.com/strimzi/client-examples.git
cd client-examples/java/kafka/producer

# step 2: build Kafka Client Example
make build
# you need scripts directory and run.sh to execute the following build because Dockerfile expects
it.
mkdir scripts
cp ../streams/scripts/run.sh scripts

# step 3: Docker build on the OpenShift
oc new-build --name=java-kafka-producer --strategy=docker --binary
oc start-build java-kafka-producer --from-dir=. --follow

# step 4: you can modify client example

# step 5: adjust image and bootstrap address in deployment yaml
sed -i "s|image: .*|image:
image-registry.openshift-image-registry.svc:5000/${NAMESPACE}/java-kafka-producer:latest|g"
../java-kafka-producer.yaml
sed -i "s|value: my-cluster-kafka-bootstrap:9092.*|value: ${KAFKA_CLUSTER}-kafka-bootstrap:9092|g"
../java-kafka-producer.yaml

# step 6: you need to also create "my-topic" and "cipot-ym" topic for Kafka streams application

# step 7: apply deployment yaml
oc apply -f ../java-kafka-producer.yaml

# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l build=java-kafka-producer
oc delete -f ../java-kafka-producer.yaml
```

Run Kafka streams client example on the OpenShift

```
# step 1: clone Kafka Client Example
git clone https://github.com/strimzi/client-examples.git
cd client-examples/java/kafka/streams

# step 2: build Kafka Client Example
make build

# step 3: Docker build on the OpenShift
oc new-build --name=java-kafka-streams --strategy=docker --binary
oc start-build java-kafka-streams --from-dir=. --follow

# step 4: you can modify client example

# step 5: adjust image and bootstrap address in deployment yaml
sed -i "s|image: .*|image:
image-registry.openshift-image-registry.svc:5000/${NAMESPACE}/java-kafka-streams:latest|g"
../java-kafka-streams.yaml
sed -i "s|value: my-source-cluster-kafka-bootstrap:9092.*|value:
${KAFKA_CLUSTER}-kafka-bootstrap:9092|g" ../java-kafka-streams.yaml

# step 6: you need to also create "my-topic" and "cipot-ym" topic for Kafka streams application

# step 7: apply deployment yaml
oc apply -f ../java-kafka-streams.yaml

# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l build=java-kafka-streams
oc delete -f ../java-kafka-streams.yaml
```

Debug

JavaVM Info

```
# Kafka broker
#   using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- jcmd kafka.Kafka VM.info

# Zookeeper
#   using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper -- jcmd
org.apache.zookeeper.server.quorum.QuorumPeerMain VM.info

# Cluster operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/strimzi-cluster-operator -- jcmd io.strimzi.operator.cluster.Main VM.info

# Topic Operator in the Entity operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- jcmd
io.strimzi.operator.topic.Main VM.info

# User Operator in the Entity operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- jcmd
```



```
io.strimzi.operator.user.Main VM.info

# Cruise Control
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control -- jcmd
com.linkedin.kafka.cruisecontrol.KafkaCruiseControlMain VM.info
```

Thread Dump

```
# Kafka broker
#   using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- jcmd kafka.Kafka Thread.print
#   using kill -3 <PID>   #[1]
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- kill -3 1

# Zookeeper
#   using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper -- jcmd
org.apache.zookeeper.server.quorum.QuorumPeerMain Thread.print
#   using kill -3 <PID>   #[1]
oc exec ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper -- kill -3 1

# Cluster operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/strimzi-cluster-operator -- jcmd io.strimzi.operator.cluster.Main Thread.print
#   using kill -3 <PID>   #[1]
oc exec deployment/strimzi-cluster-operator -- kill -3 1

# Topic Operator in the Entity operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- jcmd
io.strimzi.operator.topic.Main Thread.print
#   using kill -3 <PID>   #[1]
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- kill -3 1

# User Operator in the Entity operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- jcmd
io.strimzi.operator.user.Main Thread.print
#   using kill -3 <PID>   #[1]
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- kill -3 1

# Kafka connect
#   using kill -3 <PID>   #[1]
#   cannot know which replica to connect to. so this can be used with a single replica
oc exec deployment/${KAFKACONNECT_CLUSTER}-connect -- kill -3 1
#   or, specify pod to connect to
oc exec pod/${KAFKACONNECT_CLUSTER}-connect-XXXXXX -- kill -3 1

# Cruise Control
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control -- jcmd
com.linkedin.kafka.cruisecontrol.KafkaCruiseControlMain Thread.print
#   using kill -3 <PID>   #[1]
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control -- kill -3 1
```

[1] To get thread dump, you can also use “kill -3 <PID>”. The thread dump will be output in *java process's stdout* i.e. generally pod log, not the terminal you are using as below:

```
# For example:
```

```
# output thread dump to *java process's* stdout
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- kill -3 1
# confirm the thread dump from the pod log
oc logs ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka
```

Packet capture using tcpdump

```
# multi line commands
# step 1: specify namespace and pod name
export NAMESPACE=my-project
export POD_NAME=${KAFKA_CLUSTER}-kafka-0
# step 2: get pod id
export POD_ID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host crictl pods --namespace ${NAMESPACE} --name ${POD_NAME} -q)
# step 3: get pod process id on the node
export POD_PID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host bash -c "runc state ${POD_ID} | jq .pid")
# step 4: get packet dump using tcpdump
oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- bash -c "nsenter -n -t ${POD_PID} -- tcpdump -w - 2>/dev/null" | tee ${POD_NAME}_${date +%d_%m_%Y-%H_%M_%S-%Z}.pcap | tcpdump -r -

# single line command
# step 1: specify namespace and pod name
export NAMESPACE=my-project
export POD_NAME=${KAFKA_CLUSTER}-kafka-0
# step 2: get packet dump using tcpdump
oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- bash -c "NAMESPACE=${NAMESPACE}; POD_NAME=${POD_NAME}; ''POD_ID=$(chroot /host crictl pods --namespace ${NAMESPACE} --name ${POD_NAME} -q); POD_PID=$(chroot /host bash -c 'runc state ${POD_ID} | jq .pid'); bash -c 'nsenter -n -t ${POD_PID} -- tcpdump -l -w - 2>/dev/null'" | tee ${POD_NAME}_${date +%d_%m_%Y-%H_%M_%S-%Z}.pcap | tcpdump -r -
```

For details, refer to ["How to use tcpdump inside OpenShift v4 Pod"](#)

Heap Histogram & Statics

```
# Kafka broker
# using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- jcmd kafka.Kafka GC.class_histogram
oc exec ${KAFKA_CLUSTER}-kafka-${BROKER_ID} -c kafka -- jcmd kafka.Kafka GC.class_stats

# Zookeeper
# using jcmd since AMQ Streams 1.6
oc exec ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper -- jcmd org.apache.zookeeper.server.quorum.QuorumPeerMain GC.class_histogram
oc exec ${KAFKA_CLUSTER}-zookeeper-${ZOOKEEPER_ID} -c zookeeper -- jcmd org.apache.zookeeper.server.quorum.QuorumPeerMain GC.class_stats

# Cluster operator
# using jcmd since AMQ Streams 1.6
oc exec deployment/strimzi-cluster-operator -- jcmd io.strimzi.operator.cluster.Main GC.class_histogram
oc exec deployment/strimzi-cluster-operator -- jcmd io.strimzi.operator.cluster.Main GC.class_stats

# Topic Operator in the Entity operator
# using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- jcmd io.strimzi.operator.topic.Main GC.class_histogram
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator -- jcmd
```

```

io.strimzi.operator.topic.Main GC.class_stats

# User Operator in the Entity operator
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- jcmd
io.strimzi.operator.user.Main GC.class_histogram
oc exec deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator -- jcmd
io.strimzi.operator.user.Main GC.class_stats

# Cruise Control
#   using jcmd since AMQ Streams 1.6
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control -- jcmd
com.linkedin.kafka.cruisecontrol.KafkaCruiseControlMain GC.class_histogram
oc exec deployment/${KAFKA_CLUSTER}-cruise-control -c cruise-control -- jcmd
com.linkedin.kafka.cruisecontrol.KafkaCruiseControlMain GC.class_stats

```

Heap Dump

<https://access.redhat.com/articles/3135421> pls check this kcs for collecting the gc, thread dump, heap dump etc.

RHEL8 support tool image

```
oc run rhel8-support-tools -ti --image=registry.redhat.io/rhel8/support-tools
```

It can be used to execute commands such as curl in the namespace. And you will be the root user so you can install various packages using dnf or yum.

Change AMQ Streams Kafka image: build a image on the local machine

```

# step 1: create Dockerfile that extends amq-streams image.
vi Dockerfile
~~~
FROM registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
USER root:root
RUN microdnf install -y curl java-1.8.0-openjdk-devel
USER 1001
~~~

# step 2: build image
podman build . -t
default-route-openshift-image-registry.apps-crc.testing/${NAMESPACE}/amq-streams-debug:latest

# step 3: login and push image to the openshift repository (example of code ready container)
podman login -u kubeadmin -p $(oc whoami -t) default-route-openshift-image-registry.apps-crc.testing
podman push --tls-verify=false
default-route-openshift-image-registry.apps-crc.testing/${NAMESPACE}/amq-streams-debug:latest

# step 4: set Kafka.spec.kafka.image to the new image path.
vi kafka-my-cluster.yaml
~~~
kind: Kafka
...
spec:
  kafka:
    image: image-registry.openshift-image-registry.svc:5000/<$$NAMESPACE$$>/amq-streams-debug:latest
~~~

```

Change AMQ Streams Kafka image: build a image on the OpenShift

```
# step 1: create Dockerfile that extends amq-streams image.
vi Dockerfile
~~~
FROM registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
USER root:root
RUN microdnf install -y curl java-1.8.0-openjdk-devel
USER 1001
~~~

# step 2: Docker build on the OpenShift
oc new-build --name=amq-streams-debug --strategy=docker --binary
oc start-build amq-streams-debug --from-dir=. --follow

# step 3: set Kafka.spec.kafka.image to the new image path and apply. please adjust <$$NAMESPACE$$>
vi kafka-my-cluster.yaml
~~~
kind: Kafka
...
spec:
  kafka:
    image: image-registry.openshift-image-registry.svc:5000/<$$NAMESPACE$$>/amq-streams-debug:latest
~~~

# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l build=amq-streams-debug
```

Remote Debugging: Kafka Broker

```
# Step 1: Set KAFKA_DEBUG and JAVA_DEBUG_OPTS in the Kafka resource
~~~
Kind: Kafka
spec:
  kafka:
    template:
      kafkaContainer:
        env:
          - name: KAFKA_DEBUG
            value: "true"
          - name: JAVA_DEBUG_OPTS
            value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
~~~

# Step 2: Forward local port to the debug port on the operator pod
oc port-forward pod/${KAFKA_CLUSTER}-kafka-${BROKER_ID} 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Remote Debugging: Cluster Operator

```
# Step 1: Set JAVA_OPTS and disable health check
vi install/cluster-operator/050-Deployment-strimzi-cluster-operator.yaml
~~~
spec:
  template:
    spec:
      containers:
      - name: strimzi-cluster-operator
        env:
        - name: JAVA_OPTS
          value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
      ...
      livenessProbe:
        timeoutSeconds: 999999999
    ~~~

# Step 2: Forward local port to the debug port on the operator pod
oc port-forward deployment/strimzi-cluster-operator 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Remote Debugging: Topic Operator in Entity Operator

```
# Step 1: Set JAVA_OPTS
~~~
Kind: Kafka
spec:
  entityOperator:
    template:
      topicOperatorContainer:
        env:
        - name: JAVA_OPTS
          value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
    ~~~

# Step 2: Forward local port to the debug port on the topic operator pod
oc port-forward deployment/${KAFKA_CLUSTER}-entity-operator 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Remote Debugging: User Operator in Entity Operator

```
# Step 1: Set JAVA_OPTS
~~~
Kind: Kafka
spec:
  entityOperator:
    template:
      userOperatorContainer:
        env:
          - name: JAVA_OPTS
            value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
~~~

# Step 2: Forward local port to the debug port on the user operator pod
oc port-forward deployment/${KAFKA_CLUSTER}-entity-operator 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Remote Debugging: Kafka Connect

```
# Step 1: Set JAVA_OPTS
~~~
Kind: KafkaConnect
spec:
  template:
    connectContainer:
      env:
        - name: KAFKA_DEBUG
          value: "true"
        - name: JAVA_DEBUG_OPTS
          value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
~~~

# Step 2: Forward local port to the debug port on the kafka connect replica pod
oc port-forward pod/${KAFKACONNECT_CLUSTER}-connect-XXXXXX 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Remote Debugging: Cruise Control

```
# Step 1: Set KAFKA_OPTS
~~~
Kind: Kafka
spec:
  cruiseControl:
    template:
      cruiseControlContainer:
        env:
          - name: KAFKA_OPTS
            value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n"
~~~

# Step 2: Forward local port to the debug port on the kafka connect replica pod
oc port-forward deployment/${KAFKA_CLUSTER}-cruise-control 5005:5005

# Step 3: Debug using localhost:5005 port and IDE etc remotely.
```

Byteman: Kafka Broker

```
# step 1: create Dockerfile that extends amq-streams image. "INITIALIZE_COMMAND" will be executed
before starting Kafka.
vi Dockerfile
~~~
FROM registry.redhat.io/amq7/amq-streams-kafka-25-rhel7:1.5.0
USER root:root
RUN microdnf install -y unzip
RUN sed -i "s|# starting Kafka server.*|$(echo 'eval ${INITIALIZE_COMMAND}')|g"
/opt/kafka/kafka_run.sh
USER 1001
~~~

# step 2: Docker build on the OpenShift
oc new-build --name=amq-streams-debug --strategy=docker --binary
oc start-build amq-streams-debug --from-dir=. --follow

# step 3: set Kafka.spec.kafka.image to the new image path and apply. please adjust <${NAMESPACE}>.
and set byteman rule in "BYTEMAN_BTM".
vi kakfa-my-cluster.yaml
~~~
kind: Kafka
...
spec:
  kafka:
    image: image-registry.openshift-image-registry.svc:5000/<${NAMESPACE}>/amq-streams-debug:latest
    template:
      kafkaContainer:
        env:
          - name: INITIALIZE_COMMAND
            value: "curl -o /var/opt/kafka/byteman.zip
https://downloads.jboss.org/byteman/4.0.15/byteman-download-4.0.15-bin.zip;unzip
/var/opt/kafka/byteman.zip -d /var/opt/kafka; ls;echo \"${BYTEMAN_BTM}\" >
/var/opt/kafka/appmain.btm"
          - name: KAFKA_DEBUG
            value: "true"
          - name: JAVA_DEBUG_OPTS
            value: "-agentlib:jdwp=transport=dt_socket,address=5005,server=y,suspend=n
-javaagent:/var/opt/kafka/byteman-download-4.0.15/lib/byteman.jar=scrip: /var/opt/kafka/appmain.btm,
boot:/var/opt/kafka/byteman-download-4.0.15/lib/byteman.jar"
          - name: BYTEMAN_BTM
            value: |
              RULE rule101entry
              CLASS GroupMetadata
              METHOD completePendingTxnOffsetCommit
              AT ENTRY
              IF true
              DO delay(500);traceln("[BYTEMAN] ENTER " + $METHOD + ":" +
java.util.Arrays.asList($*));delay(500);
              ENDRULE

              RULE rule101exit
              CLASS GroupMetadata
              METHOD onTxnOffsetCommitAppend
              AT EXIT
              IF true
              DO delay(500);traceln("[BYTEMAN] EXIT " + $METHOD);delay(500);
              ENDRULE
~~~
```

```
# FYI, you can delete the build configuration and imagestream for cleanup after debugging
oc delete all -l build=amq-streams-debug
```

Tips

Dump OpenShift project

```
# For OCP 4
#   application project issue
oc adm inspect ns/myproject1 ns/myproject2 ...
#   non-application project issue
oc adm must-gather

# For OCP 3
curl -LO https://raw.githubusercontent.com/nekop/shiftbox/master/oc-dump
chmod +x ./oc-dump
./oc-dump PROJECT
```

Refer to [How to get full project information in OpenShift](#)

OpenShift events (sort by last timestamp)

```
oc get events --sort-by '.lastTimestamp'
```

To watch changes, use the “-w” option. If error occurs, try --sort-by 'metadata.creationTimestamp'

Force a pod termination once

```
oc delete pod ${KAFKA_CLUSTER}-kafka-0 --grace-period=0 --force
```

You can use this command to kill a Kafka broker or Zookeeper immediately. And OpenShift restarts the pod automatically.

Continues to force a pod termination

```
watch -n 1 oc delete pod ${KAFKA_CLUSTER}-kafka-0 --grace-period=0 --force
```

You can use this command to continue to kill a Kafka broker or Zookeeper.

I don't know if this is a good idea.

Login to Code Ready Container host node using ssh

```
ssh -i ~/.crc/machines/crc/id_rsa core@$(crc ip)
```

no matches for kind "Deployment" in version "extensions/v1beta1" error occurs

If you use old AMQ Streams 1.2 on OpenShift 4, when to install cluster operator, the following error occurs:

```
> oc apply -f install/cluster-operator/050-Deployment-strimzi-cluster-operator.yaml
error: unable to recognize "install/cluster-operator/050-Deployment-strimzi-cluster-operator.yaml":
no matches for kind "Deployment" in version "extensions/v1beta1"
```

Try the following apiVersion and spec.selector:

```
#apiVersion: extensions/v1beta1
apiVersion: apps/v1
kind: Deployment
metadata:
  name: strimzi-cluster-operator
  labels:
    app: strimzi
```



```
spec:
  replicas: 1
  selector:
    matchLabels:
      name: strimzi-cluster-operator
      strimzi.io/kind: cluster-operator
```

Long and short names for each Strimzi resource

Strimzi resource	Long name	Short name
Kafka	kafka	k
Kafka Topic	kafkatopic	kt
Kafka User	kafkauser	ku
Kafka Connect	kafkaconnect	kc
Kafka Connect S2I	kafkaconnects2i	kcs2i
Kafka Connector	kafkaconnector	kctr
Kafka Mirror Maker	kafkamirrormaker	kmm
Kafka Mirror Maker 2	kafkamirrormaker2	kmm2
Kafka Bridge	kafkabridge	kb
Kafka Rebalance	kafkarebalance	kr

Dump AMQ Streams project

Dump resources, status, and logs of AMQ Streams on OpenShift

```
./report.sh --namespace=${NAMESPACE} --cluster=${KAFKA_CLUSTER}
```

The dump includes "secrets". If you provide the dump to Red Hat support services, remove reports/secrets directory if necessary.

<https://access.redhat.com/solutions/5327571>

If you would like to get the project information itself as well, please refer to [How to get full project information in OpenShift](#)

Customizable report script

Similar to the above dump script, but the following script can be customized. Some users prefer to be clear on what commands to execute and provide the minimum information. The following script can be trimmed leaving only a few commands. And the following commands include actual Kafka topic information, not included in the above dump script.

```
#!/bin/bash
# Usage:
#   chmod +x ./reportscript.sh
#   ./reportscript.sh <KAFKA_CLUSTER> #e.g. my-cluster
#   you need to login to the namespace using AMQ Streams

# prepare
DEST=amqstreams-$(date +%Y%m%d%H%M%S).txt.gz
set -x
KAFKA_CLUSTER=$1

(
# start commands

oc version
oc whoami
oc get project $(oc project -q) -o yaml
oc status
```

```

# (using jcmd since AMQ Streams 1.6)
oc exec ${KAFKA_CLUSTER}-kafka-0 -c kafka -- jcmd kafka.Kafka VM.info
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -c zookeeper -- jcmd org.apache.zookeeper.server.quorum.QuorumPeerMain VM.info
oc exec deployment/strimzi-cluster-operator -- jcmd io.strimzi.operator.cluster.Main VM.info

oc get Kafka -o yaml
oc get KafkaTopic -o yaml
oc get KafkaConnect -o yaml
oc get KafkaConnectS2I -o yaml
oc get KafkaConnector -o yaml
oc get KafkaMirrorMaker -o yaml
oc get KafkaMirrorMaker2 -o yaml
oc get KafkaRebalance -o yaml
oc get KafkaBridge -o yaml

oc get all -l app=strimzi -o yaml

oc get all -l strimzi.io/kind -o yaml
# this command includes the following info
# oc get statefulsets ${KAFKA_CLUSTER}-kafka -o yaml
# oc get statefulsets ${KAFKA_CLUSTER}-zookeeper -o yaml
# oc get deployment/${KAFKA_CLUSTER}-entity-operator -o yaml

oc get service -o yaml
oc get pvc -o yaml -l app.kubernetes.io/name=strimzi
oc get pv -o yaml
oc get storageclasses -o yaml
oc get events --sort-by '.lastTimestamp'
oc get events --sort-by '.metadata.creationTimestamp'
oc get events

oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-configs.sh --bootstrap-server localhost:9092
--entity-type brokers --entity-name ${BROKER_ID} --describe --all

oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-features.sh --bootstrap-server localhost:9092
--describe

oc exec ${KAFKA_CLUSTER}-kafka-0 -- cat /tmp/strimzi.properties
oc exec ${KAFKA_CLUSTER}-kafka-0 -- cat /opt/kafka/custom-config/log4j.properties
oc exec ${KAFKA_CLUSTER}-kafka-0 -- ps aux
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- cat /tmp/zookeeper.properties
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- cat /opt/kafka/custom-config/log4j.properties
oc exec ${KAFKA_CLUSTER}-zookeeper-0 -- ps aux

# until AMQ Streams 1.5
oc exec ${KAFKA_CLUSTER}-kafka-0 -- bash -c "echo dump | nc localhost 2181"

oc get secret ${KAFKA_CLUSTER}-cluster-ca-cert -o 'jsonpath={.data.ca\.crt}' | base64 -d | openssl x509 -subject
-issuer -startdate -enddate -noout
oc get secret ${KAFKA_CLUSTER}-clients-ca-cert -o 'jsonpath={.data.ca\.crt}' | base64 -d | openssl x509 -subject
-issuer -startdate -enddate -noout

oc logs ${KAFKA_CLUSTER}-kafka-0 -c kafka
oc logs -p ${KAFKA_CLUSTER}-kafka-0 -c kafka
oc logs ${KAFKA_CLUSTER}-kafka-0 -c tls-sidecar
oc logs -p ${KAFKA_CLUSTER}-kafka-0 -c tls-sidecar

oc logs ${KAFKA_CLUSTER}-zookeeper-0 -c zookeeper
oc logs -p ${KAFKA_CLUSTER}-zookeeper-0 -c zookeeper
oc logs ${KAFKA_CLUSTER}-zookeeper-0 -c tls-sidecar
oc logs -p ${KAFKA_CLUSTER}-zookeeper-0 -c tls-sidecar

oc logs deployment/strimzi-cluster-operator
oc logs -p deployment/strimzi-cluster-operator
oc logs deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator
oc logs -p deployment/${KAFKA_CLUSTER}-entity-operator -c topic-operator
oc logs deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator

```

```
oc logs -p deployment/${KAFKA_CLUSTER}-entity-operator -c user-operator

TOPICS=$(oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - KAFKA_JVM_PERFORMANCE_OPTS=-XX:ParallelGCThreads=1
bin/kafka-topics.sh --zookeeper localhost:2181 --list)
for TOPIC in $TOPICS
do
    oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-topics.sh --zookeeper localhost:2181
--describe --topic $TOPIC
done

oc exec -it ${KAFKA_CLUSTER}-kafka-0 -c kafka -- env - bin/kafka-consumer-groups.sh --bootstrap-server
localhost:9092 --describe --all-groups

) 2>&1 | gzip > $DEST
echo "Generated $DEST"
```

You can split a dump file to files per command using the following command:

```
zcat *.txt.gz | awk 'BEGIN{f=""} match($0, /^+ ((oc|date).*)$/, a){f=a[1] ".txt"; gsub(/[
\\=]/,"_",f);} {print $0 >> f}'
```

References

- Red Hat AMQ Documentation
 - https://access.redhat.com/documentation/en-us/red_hat_amq/2020.q4/
- Red Hat AMQ 7 Supported Configuration
 - <https://access.redhat.com/articles/2791941>
- Red Hat AMQ 7 Component Details
 - <https://access.redhat.com/articles/3188232>
- Red Hat AMQ Streams Download
 - <https://access.redhat.com/jbossnetwork/restricted/listSoftware.html?downloadType=distributions&product=jboss.amq.streams>
- Kafka Documentation
 - <https://kafka.apache.org/documentation/>
- Strimzi Documentation
 - <https://strimzi.io/docs/operators/latest/using.html>
- AMQ Streams sizing recommendations
 - <https://docs.google.com/document/d/1P02sMeLLGylcfec6TvdLkIXPwTRE61hwRE9nREHbSA/edit#>
 - https://docs.google.com/presentation/d/1GEIvsWu9F9MJnPVzs_fpuzmyPRI2-6bUH3aAt_rP1A/edit#slide=id.g9c3ac4421_0_58
- Strimzi 2.4 > 10.6. Tuning Kafka configuration
 - <https://strimzi.io/docs/operators/latest/using.html#assembly-tuning-config-str>
- AMQ Streams - FAQ
 - <https://docs.google.com/document/d/1pEUSHl4IFcGqXAUiu98MtowdAmH5V4qO60vDq0-ewSw/edit?usp=sharing>
- Red Hat Integration Capacity Guidance
 - <https://docs.google.com/document/d/1CumcxzS5ykmrsGEwnaL4FaNxYx2lBag1ZNNJQ6hwmjQ/edit#heading=h.mn40jvoqksnc>
- AMQ Streams Enablement Training Agenda
 - https://docs.google.com/spreadsheets/d/17Et8LEo7SLsxjqtEgrlyzCqjBK1sH6dizVBo_W6rDJQ/edit#gid=1411071195
- Internal Training Sessions for Fuse and AMQ (Integration)
 - https://source.redhat.com/groups/public/gss-fusesource/fuseeducationcenter/fuse_education_center_wiki/internal_training_sessions_for_fuse_and_amq_integration#jive_content_id_AMQStreams_Kafka
- AMQ Streams on OpenShift staging documentation (master branch)

- <https://gitlab.cee.redhat.com/AMQ7-documentation/amq-streams/-/tree/master/books-ocp>
- AMQ Streams on Red Hat Enterprise Linux staging documentation (master branch)
 - <https://gitlab.cee.redhat.com/AMQ7-documentation/amq-streams/-/tree/master/books-rhel>

Unorganized

Login to the pod's node using debug container

```
# step 1: specify namespace and pod name
NAMESPACE=my-project
POD_NAME=${KAFKA_CLUSTER}-kafka-0
# step 2: get pod id
POD_ID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host
crictl pods --namespace ${NAMESPACE} --name ${POD_NAME} -q)
# step 3: get pod process id on the node
POD_PID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host
bash -c "runc state ${POD_ID} | jq .pid")
# step 4: login the node
oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}')
```

Execute GDB from a node

```
# step 1: login to the pod's node using debug container
#   refer to the above 'login to the pod's node using debug container' section
# step 2: install GDB
dnf install -y gdb
# step 3: execute GDB
gdb -p $PID
```

Kill a connection using GDB

```
# step 1: login to the pod's node using debug container
#   refer to the above 'login to the pod's node using debug container' section
# step 2: install GDB
dnf install -y gdb
```

```
# step 3: enter to pod's network namespace
nsenter -n -t <POD_PID>
# step 4: check the port/pid mapping
netstat -antp
# step 5: check the pid/file descriptor mapping: It has postfix of "u" like "120u" which means "120".
lsof -np <pid>
# step 6: close the connection's file descriptor using GDB
gdb -p $PID --batch -ex '{call (int)close(<file descriptor id>)}'
```

```
export POD_ID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host crictl pods --namespace ${NAMESPACE} --name ${POD_NAME} -q)
export POD_PID=$(oc debug node/$(oc get pod/${POD_NAME} -o=jsonpath='{.spec.nodeName}') -- chroot /host bash -c "runc state ${POD_ID} | jq .pid")
env | grep POD
```

```
nsenter -n -t XXXXX
chroot /host
iptables -I INPUT -p tcp --dport 21810 -j REJECT --reject-with tcp-reset
iptables -I OUTPUT -p tcp --sport 21810 -j REJECT --reject-with tcp-reset

iptables -D INPUT -p tcp --dport 21810 -j REJECT --reject-with tcp-reset
iptables -D OUTPUT -p tcp --sport 21810 -j REJECT --reject-with tcp-reset
```

<https://access.redhat.com/solutions/30303>

Show actual network policy

```
oc get NetworkPolicy/${KAFKA_CLUSTER}-network-policy-kafka -o yaml
oc get NetworkPolicy/${KAFKA_CLUSTER}-network-policy-zookeeper -o yaml
```